

# MP4

April 17, 2022

## 1 Generative Adversarial Networks

For this part of the assignment you implement two different types of generative adversarial networks. We will train the networks on a dataset of cat face images.

```
[1]: import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder
import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2
```

```
[2]: from gan.train import train
```

```
[3]: device = torch.device("cuda:0" if torch.cuda.is_available() else "gpu")
```

## 2 GAN loss functions

In this assignment you will implement two different types of GAN cost functions. You will first implement the loss from the [original GAN paper](#). You will also implement the loss from [LS-GAN](#).

### 2.0.1 GAN loss

**TODO:** Implement the `discriminator_loss` and `generator_loss` functions in `gan/losses.py`.

The generator loss is given by:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log (1 - D(G(z)))]$$

Note that these are negated from the equations presented earlier as we will be *minimizing* these losses.

**HINTS:** You should use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss. The BCE loss is needed to compute the log probability of the true label given the logits output from the discriminator. Given a score  $s \in \mathbb{R}$  and a label  $y \in \{0, 1\}$ , the binary cross entropy loss is

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

Instead of computing the expectation of  $\log D(G(z))$ ,  $\log D(x)$  and  $\log(1 - D(G(z)))$ , we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing.

```
[4]: from gan.losses import discriminator_loss, generator_loss
```

## 2.0.2 Least Squares GAN loss

**TODO:** Implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

**HINTS:** Instead of computing the expectation, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing. When plugging in for  $D(x)$  and  $D(G(z))$  use the direct output from the discriminator (`scores_real` and `scores_fake`).

```
[16]: from gan.losses import ls_discriminator_loss, ls_generator_loss
```

## 3 GAN model architecture

**TODO:** Implement the Discriminator and Generator networks in `gan/models.py`.

We recommend the following architectures which are inspired by [DCGAN](#):

### Discriminator:

- convolutional layer with `in_channels=3, out_channels=128, kernel=4, stride=2`
- convolutional layer with `in_channels=128, out_channels=256, kernel=4, stride=2`
- batch norm

- convolutional layer with in\_channels=256, out\_channels=512, kernel=4, stride=2
- batch norm
- convolutional layer with in\_channels=512, out\_channels=1024, kernel=4, stride=2
- batch norm
- convolutional layer with in\_channels=1024, out\_channels=1, kernel=4, stride=1

Use padding = 1 (not 0) for all the convolutional layers.

Instead of Relu we LeakyReLU throughout the discriminator (we use a negative slope value of 0.2). You can use simply use relu as well.

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

### Generator:

**Note:** In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution). This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with in\_channels=NOISE\_DIM, out\_channels=1024, kernel=4, stride=1
- batch norm
- transpose convolution with in\_channels=1024, out\_channels=512, kernel=4, stride=2
- batch norm
- transpose convolution with in\_channels=512, out\_channels=256, kernel=4, stride=2
- batch norm
- transpose convolution with in\_channels=256, out\_channels=128, kernel=4, stride=2
- batch norm
- transpose convolution with in\_channels=128, out\_channels=3, kernel=4, stride=2

The output of the final layer of the generator network should have a `tanh` nonlinearity to output values between -1 and 1. The output should be a 3x64x64 tensor for each sample (equal dimensions to the images from the dataset).

```
[6]: from gan.models import Discriminator, Generator
```

## 4 Data loading

The cat images we provide are RGB images with a resolution of 64x64. In order to prevent our discriminator from overfitting, we will need to perform some data augmentation.

**TODO:** Implement data augmentation by adding new transforms to the cell below. At the minimum, you should have a RandomCrop and a ColorJitter, but we encourage you to experiment with different augmentations to see how the performance of the GAN changes. See <https://pytorch.org/vision/stable/transforms.html>.

```
[7]: batch_size = 32
imsize = 64
cat_root = './cats'
```

```

cat_train = ImageFolder(root=cat_root, transform=transforms.Compose([
    transforms.ToTensor(),

    # Example use of RandomCrop:
    transforms.Resize(int(1.15 * imsize)),
    transforms.RandomCrop(imsize),
])

cat_loader_train = DataLoader(cat_train, batch_size=batch_size, drop_last=True)

```

#### 4.0.1 Visualize dataset

```
[8]: from gan.utils import show_images

imgs = cat_loader_train.__iter__().next()[0].numpy().squeeze()
show_images(imgs, color=True)
```



## 5 Training

**TODO:** Fill in the training loop in `gan/train.py`.

```
[9]: NOISE_DIM = 100  
NUM_EPOCHS = 50  
learning_rate = 0.001
```

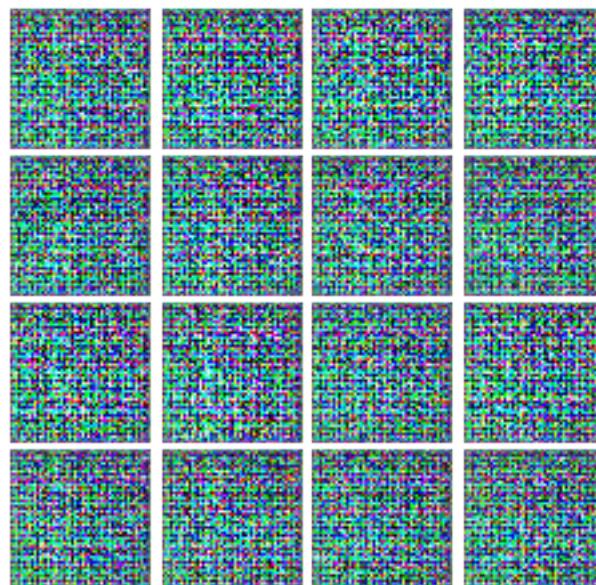
### 5.0.1 Train GAN

```
[10]: D = Discriminator().to(device)  
G = Generator(noise_dim=NOISE_DIM).to(device)
```

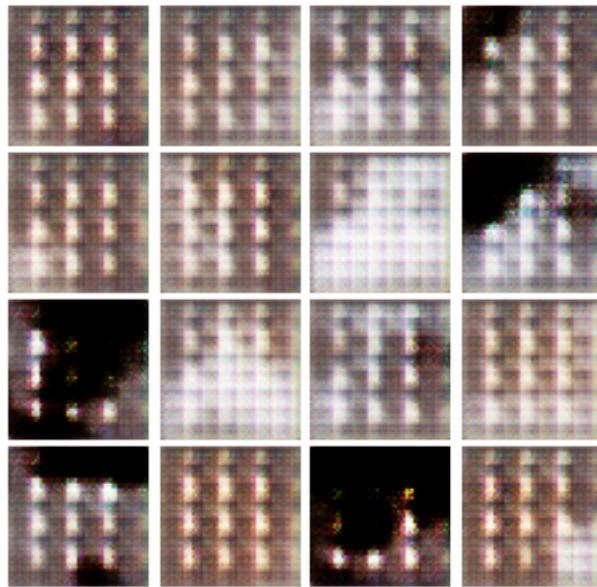
```
[11]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))  
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))
```

```
[12]: # original gan  
train(D, G, D_optimizer, G_optimizer, discriminator_loss,  
      generator_loss, num_epochs=NUM_EPOCHS, show_every=250,  
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

EPOCH: 1  
Iter: 0, D: 0.5419, G:0.2016

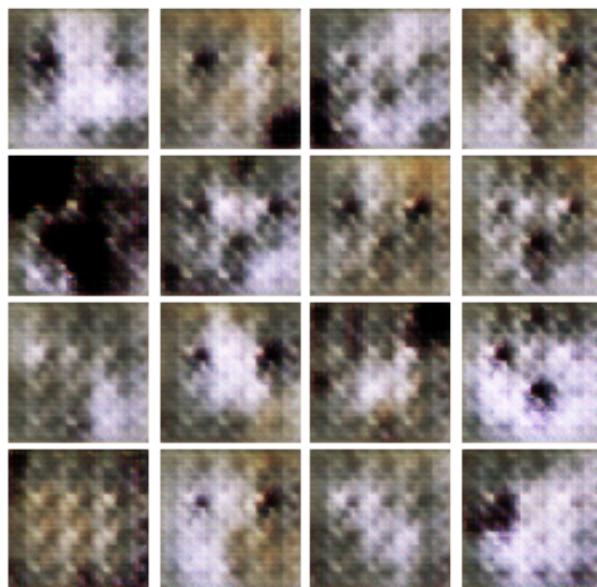


Iter: 250, D: 0.5134, G:1.418

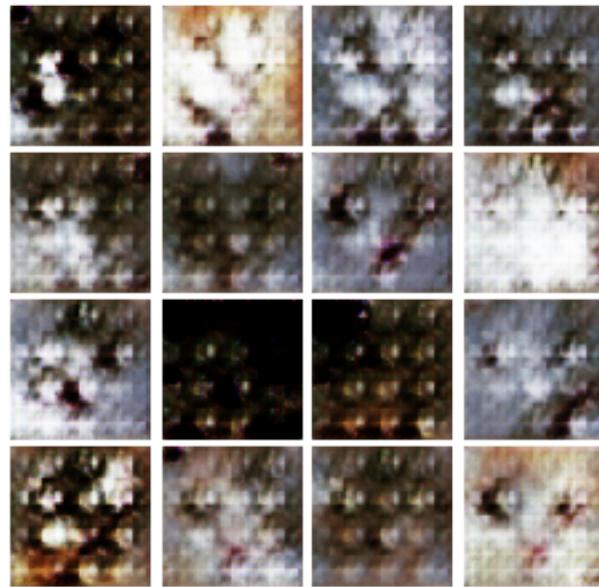


EPOCH: 2

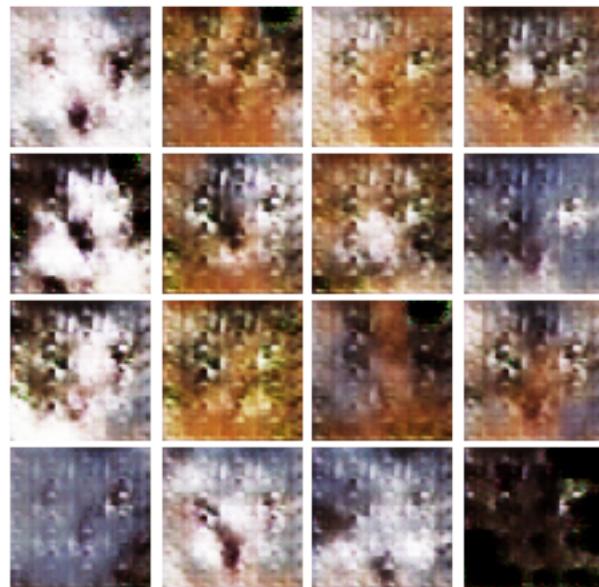
Iter: 500, D: 0.2871, G:0.997



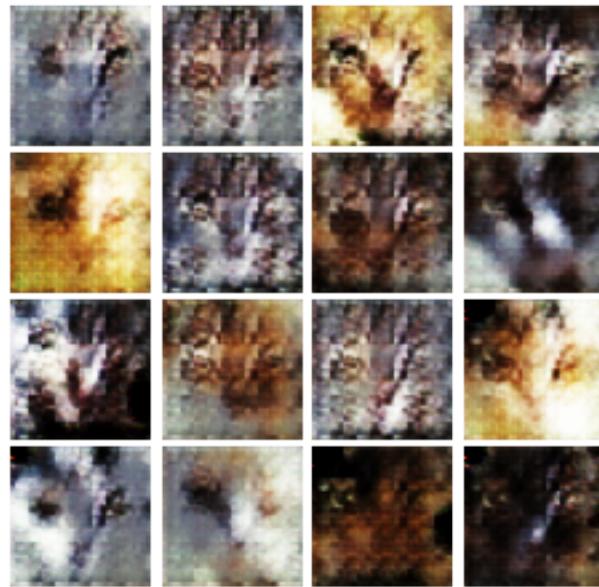
Iter: 750, D: 0.2497, G:1.194



EPOCH: 3  
Iter: 1000, D: 0.7794, G:0.9906



Iter: 1250, D: 0.279, G:1.692



EPOCH: 4  
Iter: 1500, D: 0.6582, G:3.069

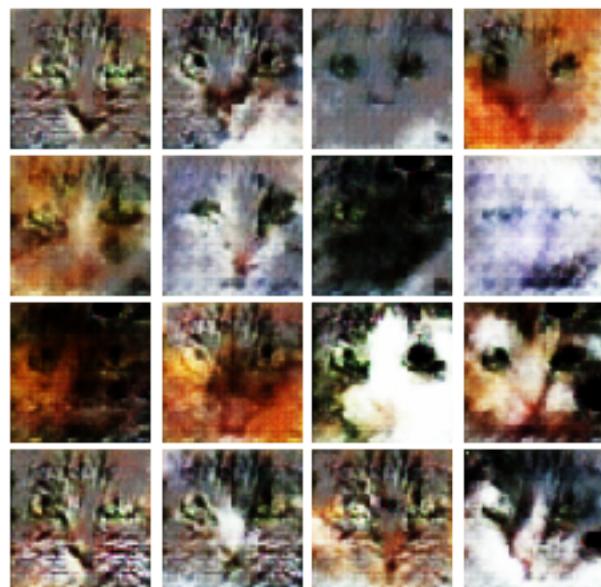


Iter: 1750, D: 0.2601, G:0.8415



EPOCH: 5

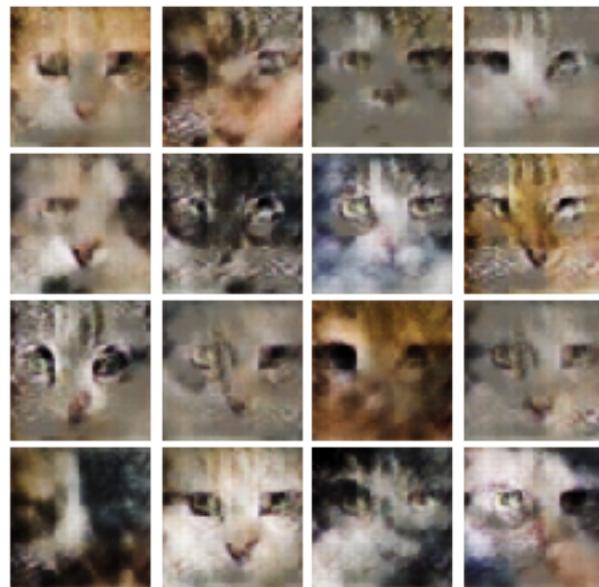
Iter: 2000, D: 0.3582, G:2.266



Iter: 2250, D: 0.2428, G:2.412



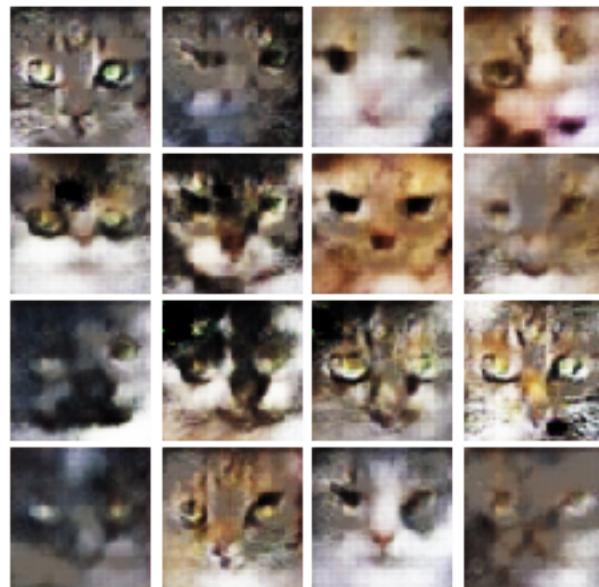
EPOCH: 6  
Iter: 2500, D: 0.3178, G:0.7039



Iter: 2750, D: 0.1657, G:2.017



EPOCH: 7  
Iter: 3000, D: 0.3682, G:4.999



Iter: 3250, D: 0.01325, G:3.828



EPOCH: 8

Iter: 3500, D: 0.185, G:2.988



Iter: 3750, D: 0.06524, G:3.066

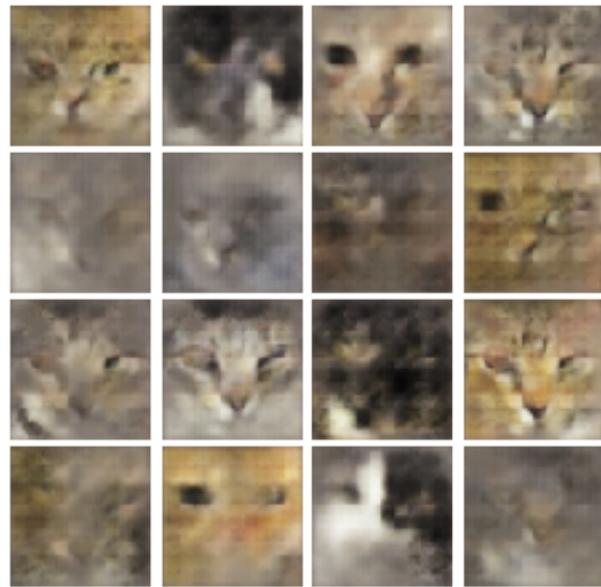


EPOCH: 9

Iter: 4000, D: 0.1011, G:1.772

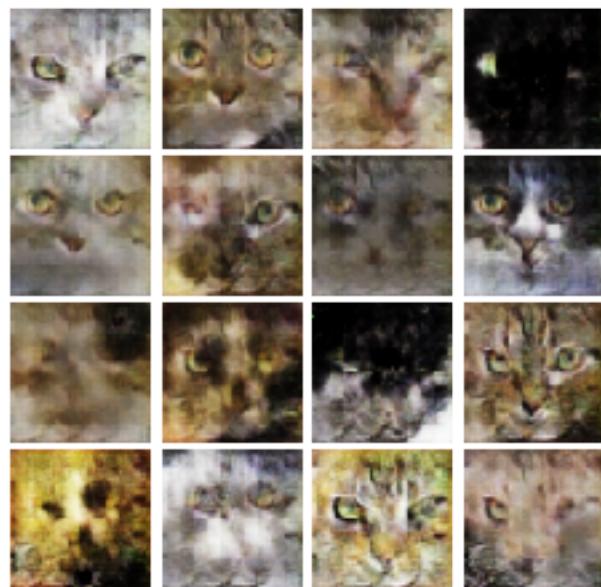


Iter: 4250, D: 0.09664, G:2.508



EPOCH: 10

Iter: 4500, D: 0.07991, G:3.642



Iter: 4750, D: 0.03636, G:3.717



EPOCH: 11  
Iter: 5000, D: 0.246, G:8.392



Iter: 5250, D: 0.02255, G:4.29



EPOCH: 12

Iter: 5500, D: 0.1662, G:10.71



Iter: 5750, D: 0.02816, G:4.73

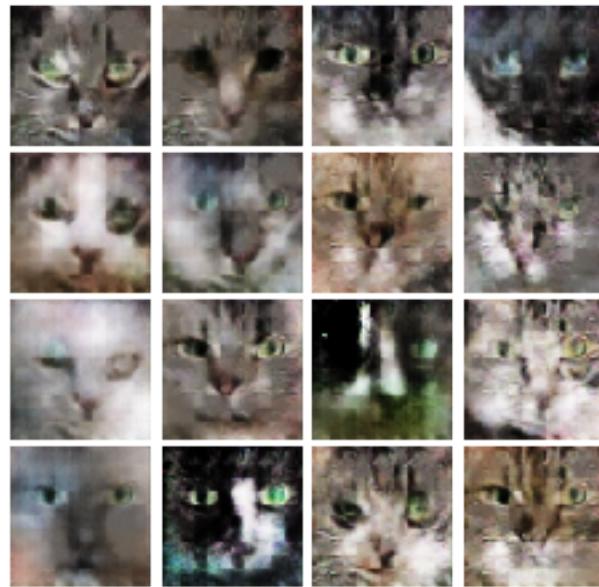


EPOCH: 13

Iter: 6000, D: 0.04154, G: 4.824



Iter: 6250, D: 0.0416, G: 4.637



EPOCH: 14

Iter: 6500, D: 0.01875, G: 4.207



Iter: 6750, D: 0.08856, G: 3.825



EPOCH: 15

Iter: 7000, D: 0.01533, G:2.932



Iter: 7250, D: 0.1214, G:4.029



EPOCH: 16

Iter: 7500, D: 0.1817, G:7.146



Iter: 7750, D: 0.06882, G:3.563



EPOCH: 17

Iter: 8000, D: 0.04477, G:3.546



Iter: 8250, D: 0.09635, G:5.121

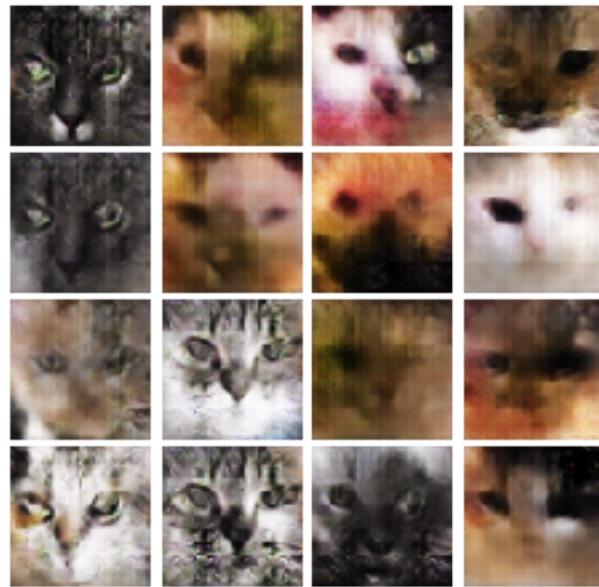


EPOCH: 18

Iter: 8500, D: 0.07346, G: 4.184



Iter: 8750, D: 0.03715, G: 2.944



EPOCH: 19

Iter: 9000, D: 0.07545, G:4.252

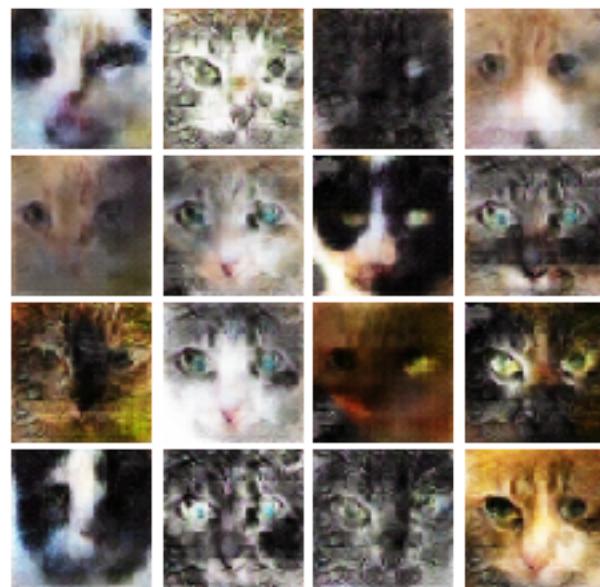


Iter: 9250, D: 0.04984, G:3.496



EPOCH: 20

Iter: 9500, D: 0.03367, G: 4.386



Iter: 9750, D: 0.008548, G: 5.921

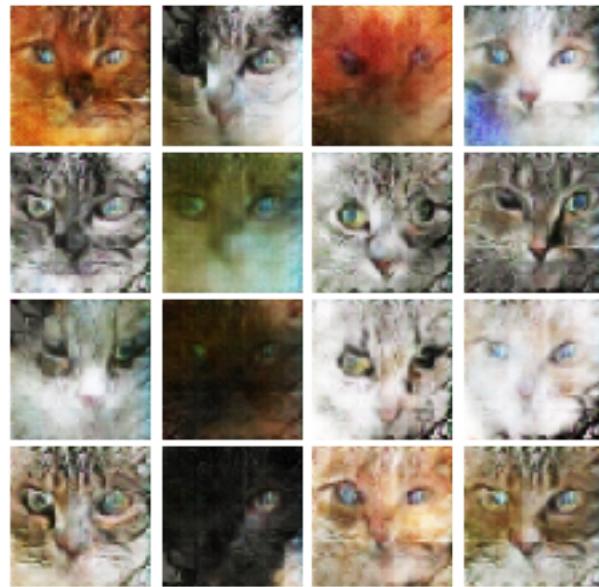


EPOCH: 21

Iter: 10000, D: 0.072, G:3.505



Iter: 10250, D: 0.03173, G:3.53



EPOCH: 22

Iter: 10500, D: 0.04577, G:4.67

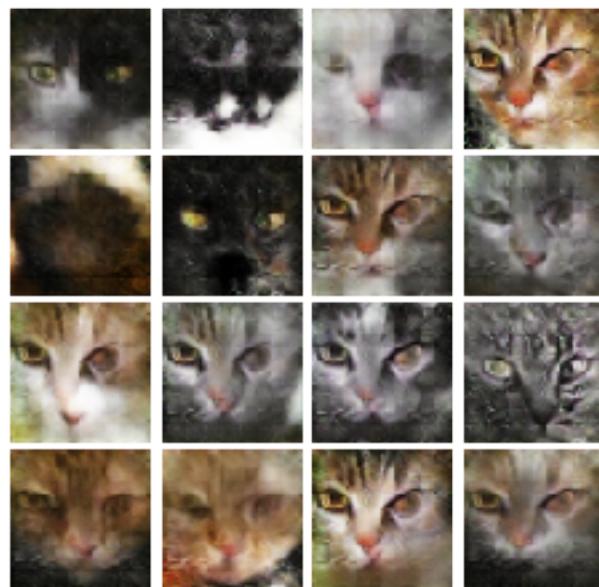


Iter: 10750, D: 0.02583, G:3.049



EPOCH: 23

Iter: 11000, D: 0.03306, G:4.18



Iter: 11250, D: 0.2814, G:6.919



EPOCH: 24

Iter: 11500, D: 0.03427, G:3.726



Iter: 11750, D: 0.08617, G:6.449

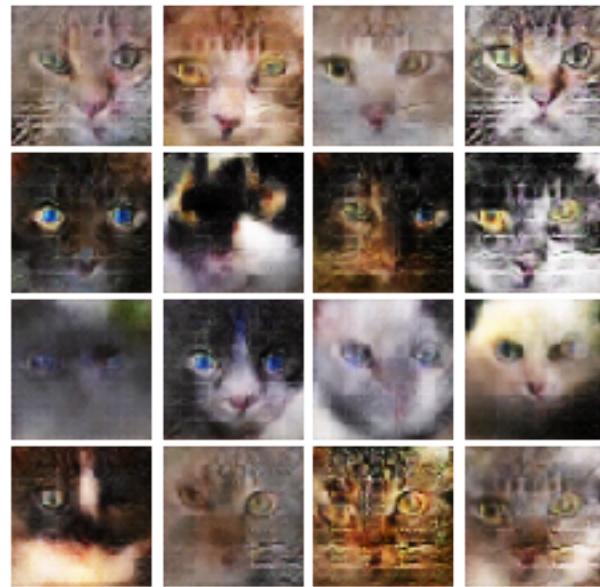


EPOCH: 25

Iter: 12000, D: 0.01915, G:3.488

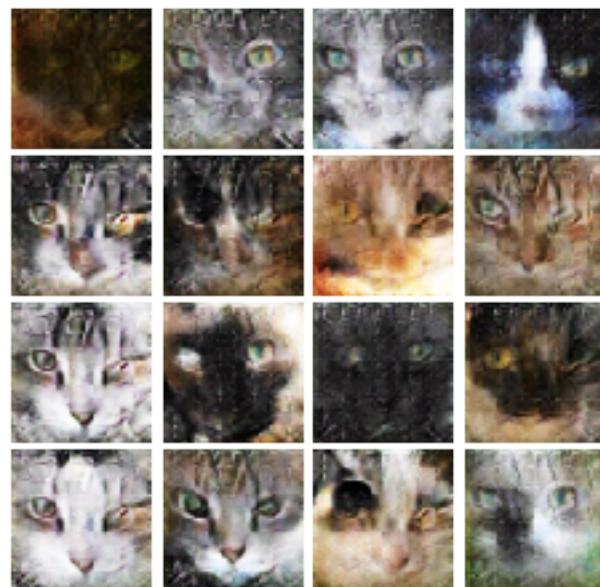


Iter: 12250, D: 0.02077, G:4.9



EPOCH: 26

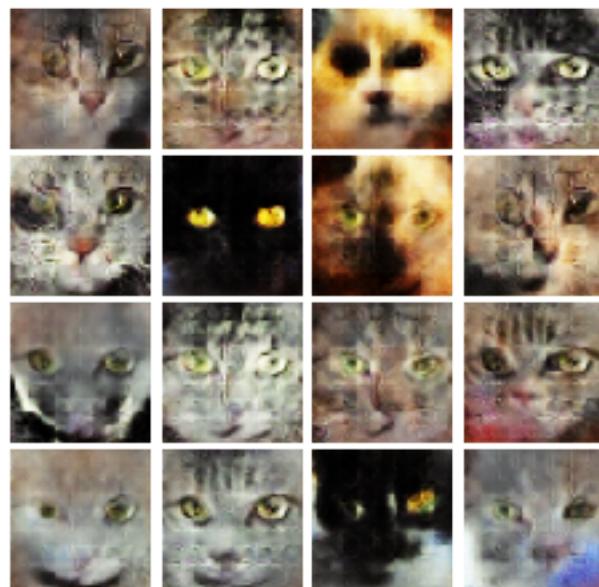
Iter: 12500, D: 0.09245, G:2.249



Iter: 12750, D: 0.06511, G:3.141



EPOCH: 27  
Iter: 13000, D: 0.05285, G:4.435



Iter: 13250, D: 0.02902, G:6.772



EPOCH: 28

Iter: 13500, D: 0.01659, G:3.788

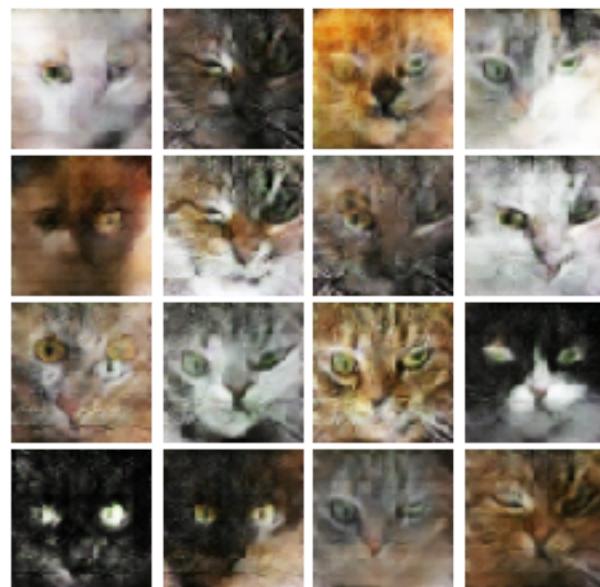


Iter: 13750, D: 0.03198, G:5.224

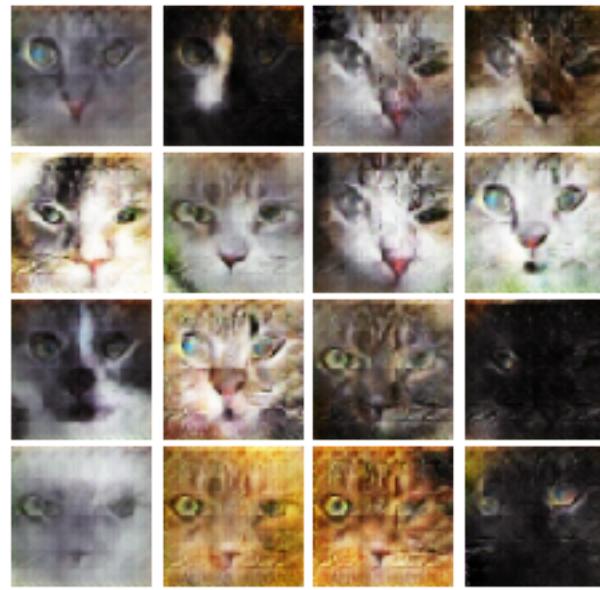


EPOCH: 29

Iter: 14000, D: 0.1571, G: 5.411

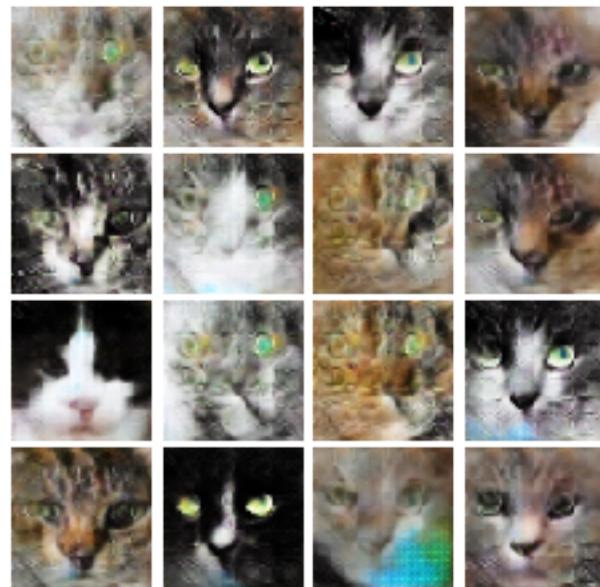


Iter: 14250, D: 0.03265, G: 4.094

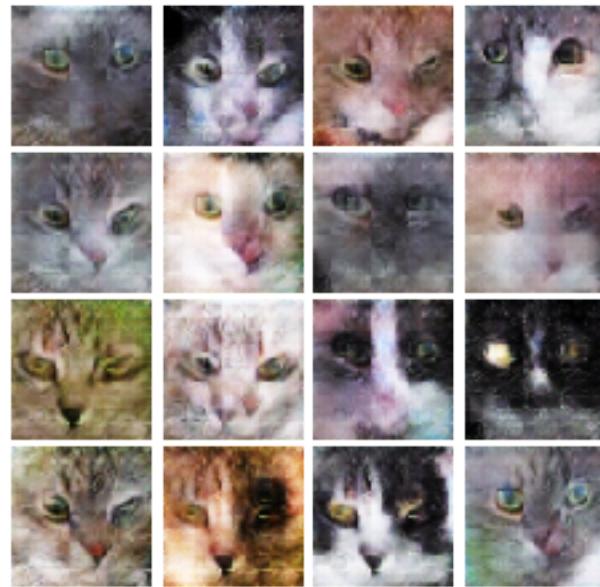


EPOCH: 30

Iter: 14500, D: 0.01531, G:4.302



Iter: 14750, D: 0.02788, G:3.681



EPOCH: 31

Iter: 15000, D: 0.03224, G:4.691



Iter: 15250, D: 0.006476, G:5.744



EPOCH: 32

Iter: 15500, D: 0.02498, G:5.908



EPOCH: 33

Iter: 15750, D: 0.02207, G:4.454



Iter: 16000, D: 0.1395, G:7.691



EPOCH: 34  
Iter: 16250, D: 0.02109, G:4.569



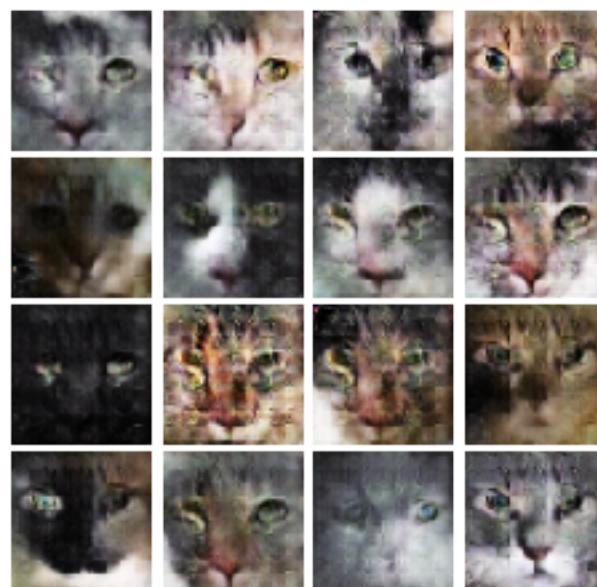
Iter: 16500, D: 0.02601, G:6.723



EPOCH: 35  
Iter: 16750, D: 0.0222, G:4.822

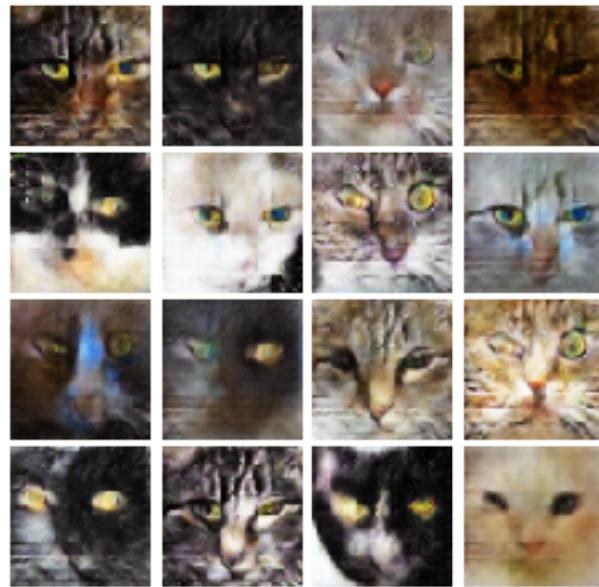


Iter: 17000, D: 0.01012, G:6.653

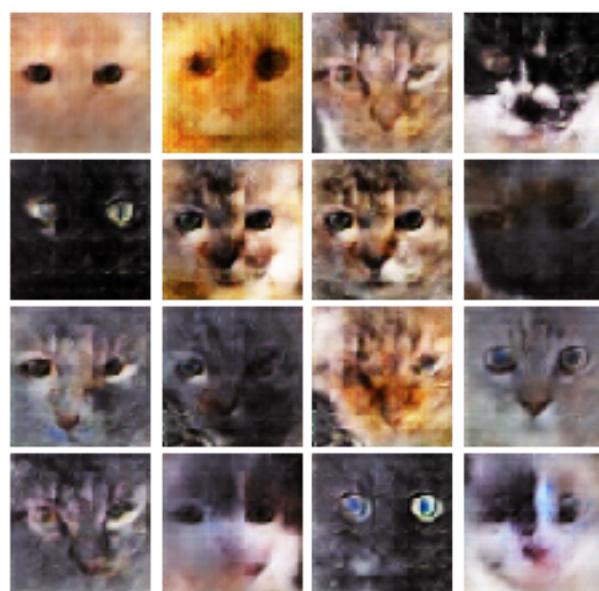


EPOCH: 36

Iter: 17250, D: 0.02458, G:3.531



Iter: 17500, D: 0.1081, G:2.549



EPOCH: 37  
Iter: 17750, D: 0.02056, G:5.022



Iter: 18000, D: 0.05432, G:6.221

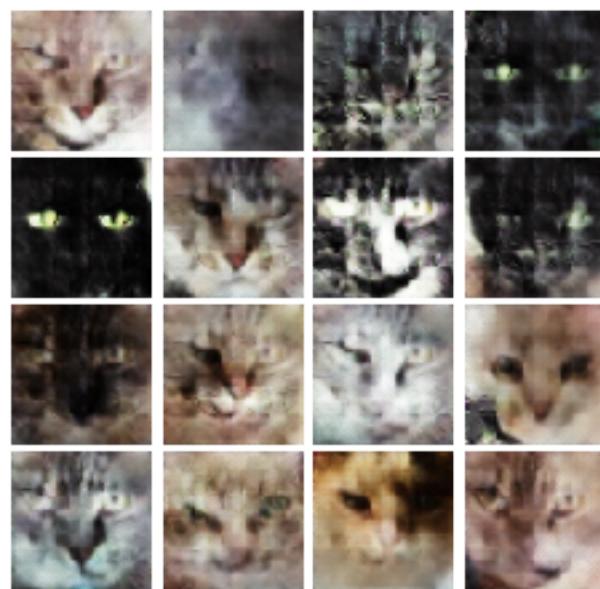


EPOCH: 38

Iter: 18250, D: 0.02272, G:4.345



Iter: 18500, D: 0.03192, G:6.023



EPOCH: 39  
Iter: 18750, D: 0.02688, G:3.838



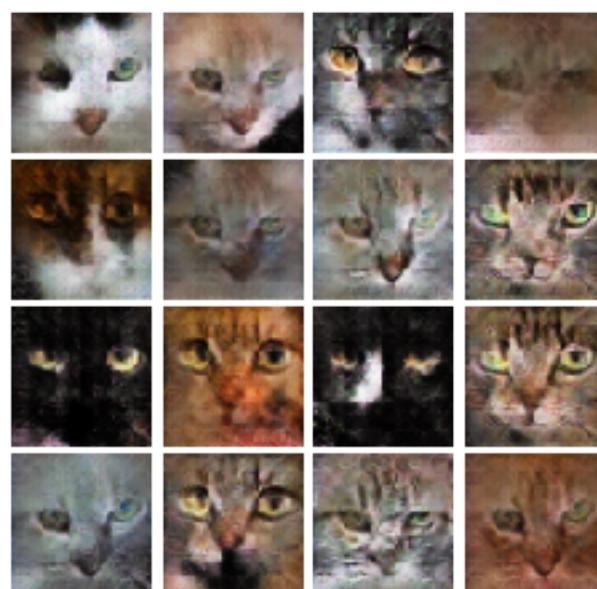
Iter: 19000, D: 0.01843, G:4.612



EPOCH: 40  
Iter: 19250, D: 0.01598, G:3.572



Iter: 19500, D: 0.05785, G:3.486



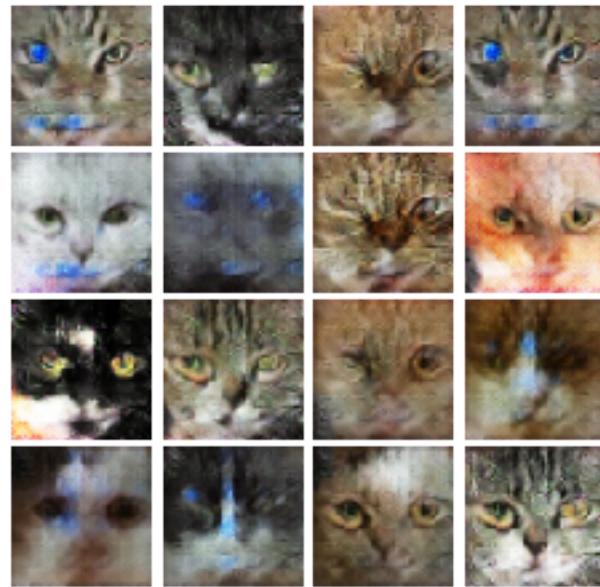
EPOCH: 41  
Iter: 19750, D: 0.01231, G:4.04



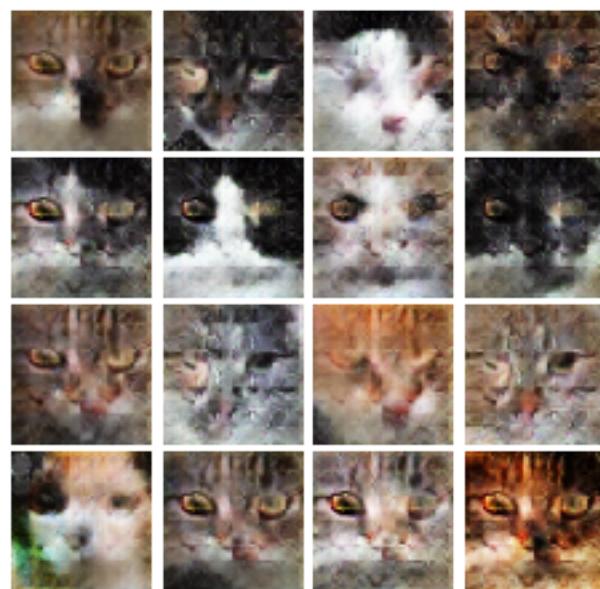
Iter: 20000, D: 0.008936, G:2.982



EPOCH: 42  
Iter: 20250, D: 0.03987, G:4.177



Iter: 20500, D: 0.00758, G:4.233



EPOCH: 43  
Iter: 20750, D: 0.06202, G:2.718

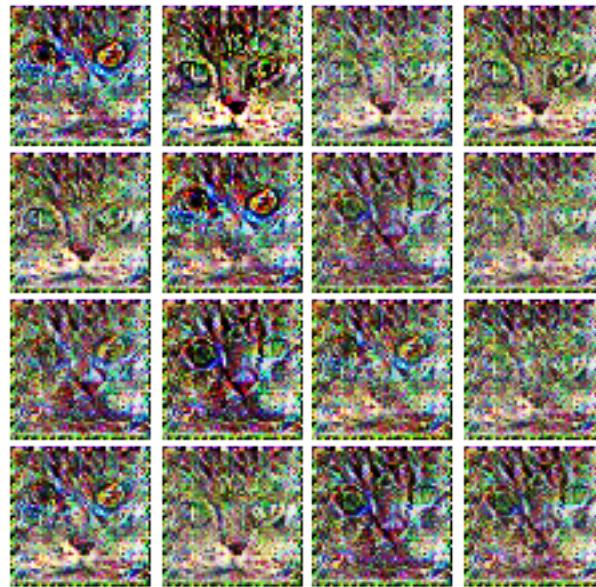


Iter: 21000, D: 0.003143, G:7.974



EPOCH: 44

Iter: 21250, D: 0.0006643, G:7.217



Iter: 21500, D: 0.01106, G:5.186

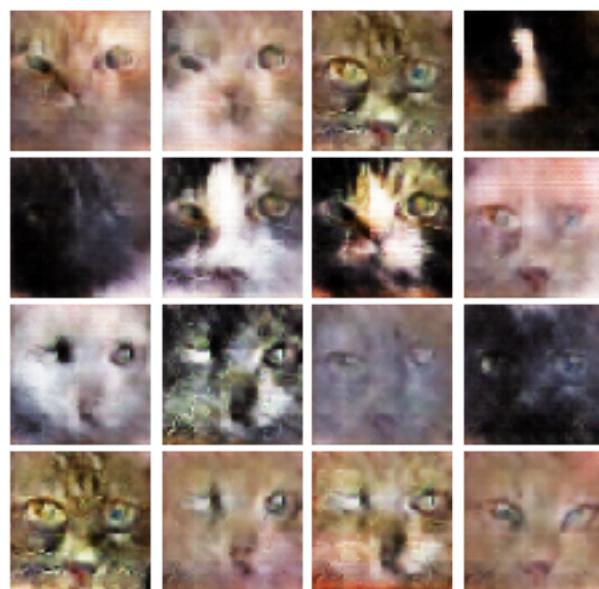


EPOCH: 45

Iter: 21750, D: 0.01132, G:4.39



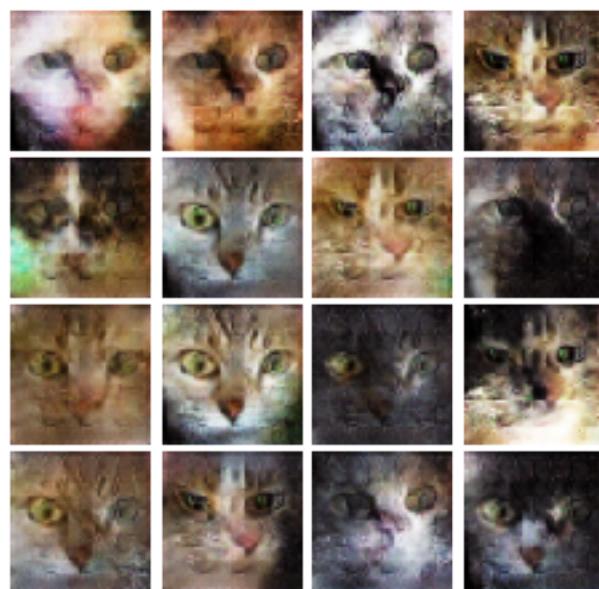
Iter: 22000, D: 0.007328, G:4.708



EPOCH: 46  
Iter: 22250, D: 0.0142, G:11.16



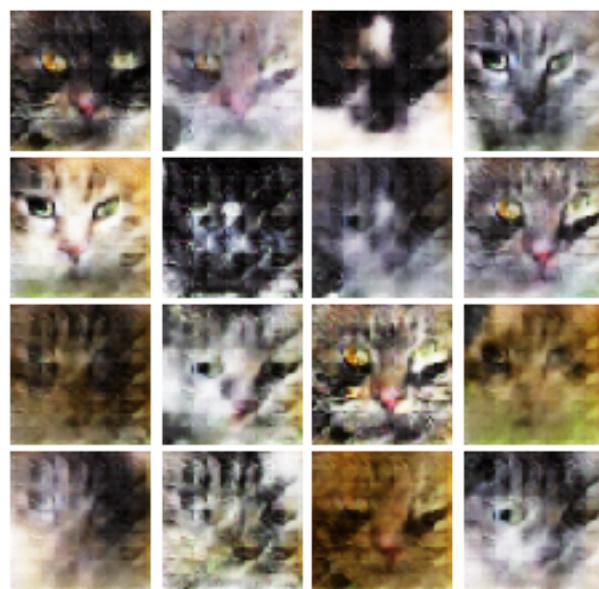
Iter: 22500, D: 0.03819, G:6.969



EPOCH: 47  
Iter: 22750, D: 0.02217, G:4.886



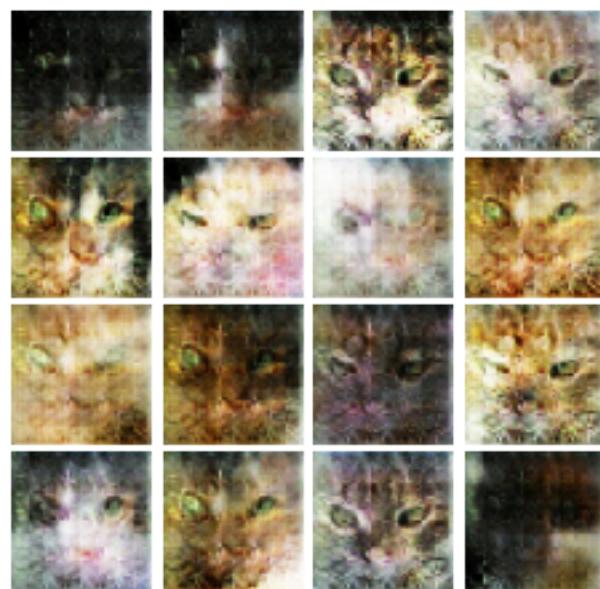
Iter: 23000, D: 0.01085, G:6.207



EPOCH: 48  
Iter: 23250, D: 0.008789, G:5.004



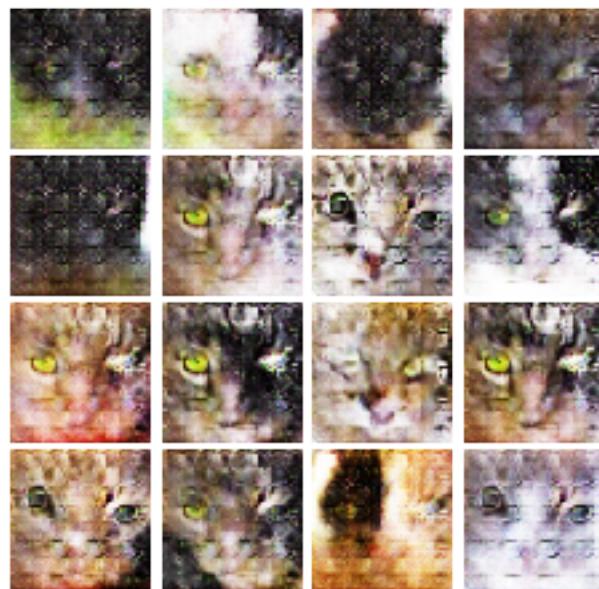
Iter: 23500, D: 0.0534, G:6.307



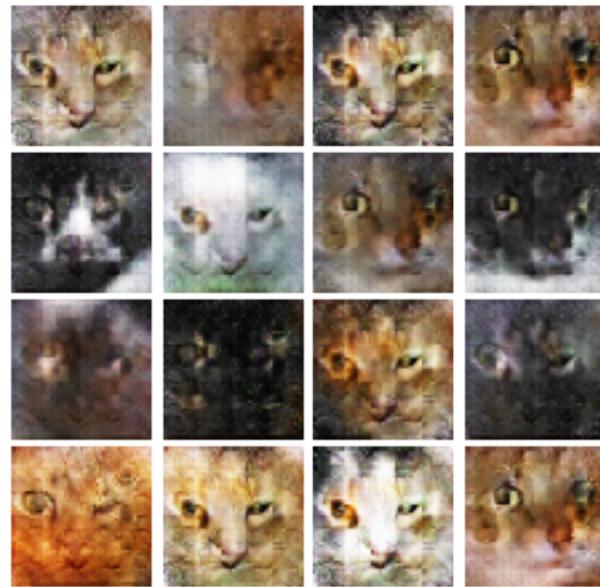
EPOCH: 49  
Iter: 23750, D: 0.08832, G:5.105



Iter: 24000, D: 0.0116, G:6.107



EPOCH: 50  
Iter: 24250, D: 0.01134, G:6.116



Iter: 24500, D: 0.02577, G: 4.939



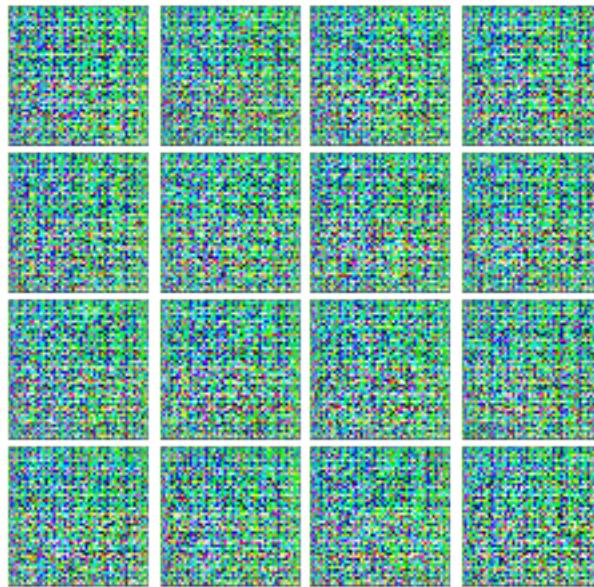
## 5.0.2 Train LS-GAN

```
[13]: D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)

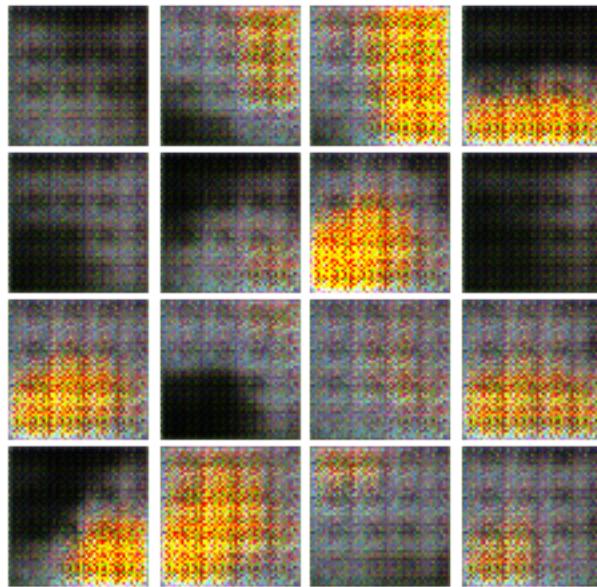
[14]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

[17]: # ls-gan
train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
      ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

EPOCH: 1  
Iter: 0, D: 0.9445, G:255.7

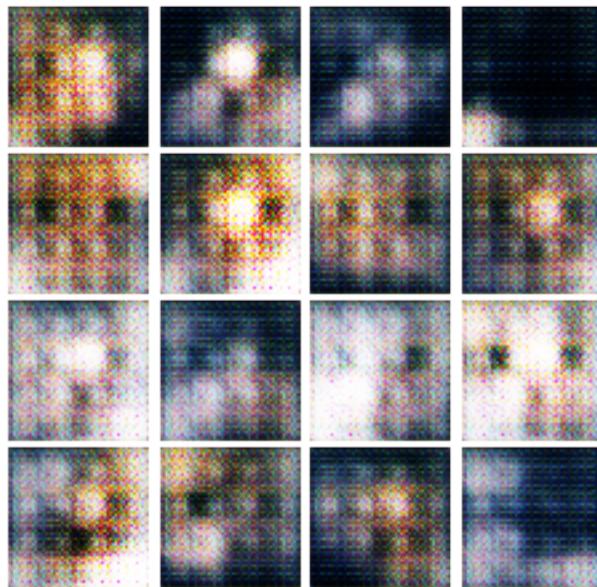


Iter: 250, D: 0.1592, G:0.8952

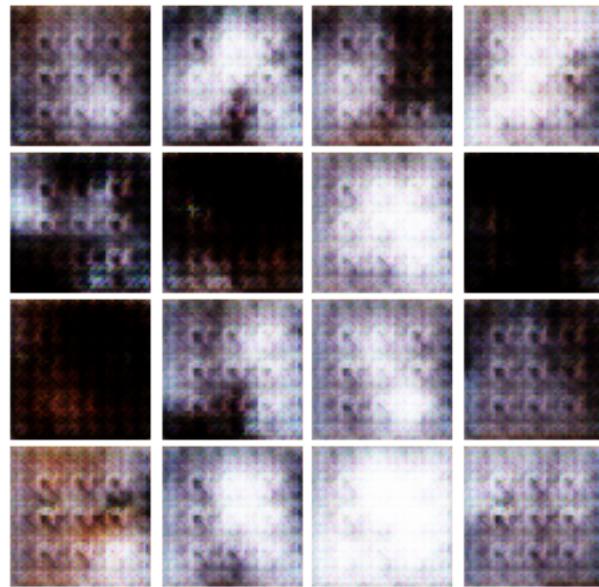


EPOCH: 2

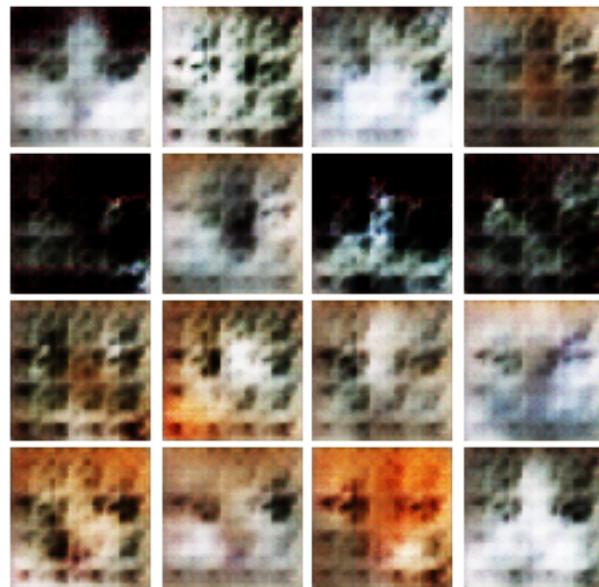
Iter: 500, D: 0.1965, G: 0.9881



Iter: 750, D: 0.0825, G: 1.38



EPOCH: 3  
Iter: 1000, D: 0.178, G:0.5018



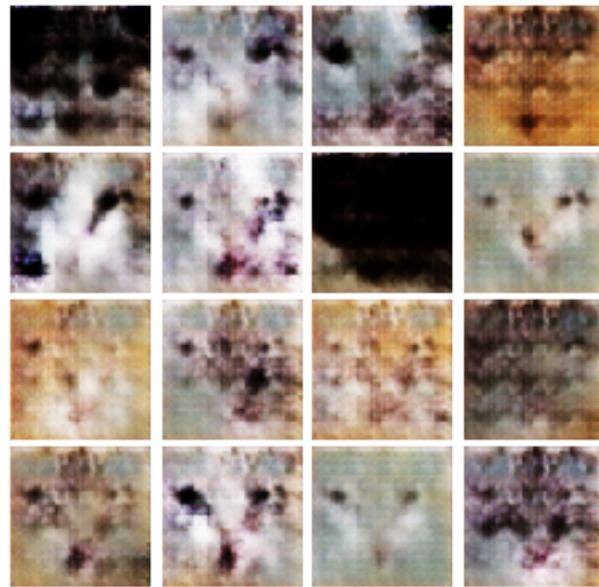
Iter: 1250, D: 0.1428, G:0.5411



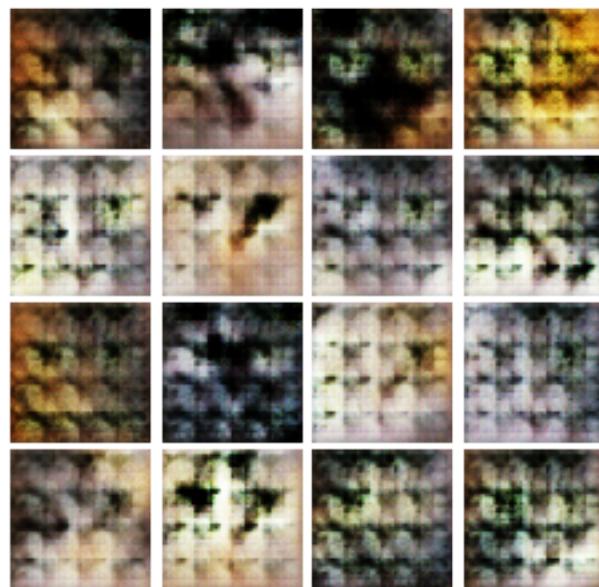
EPOCH: 4  
Iter: 1500, D: 0.2193, G: 0.705



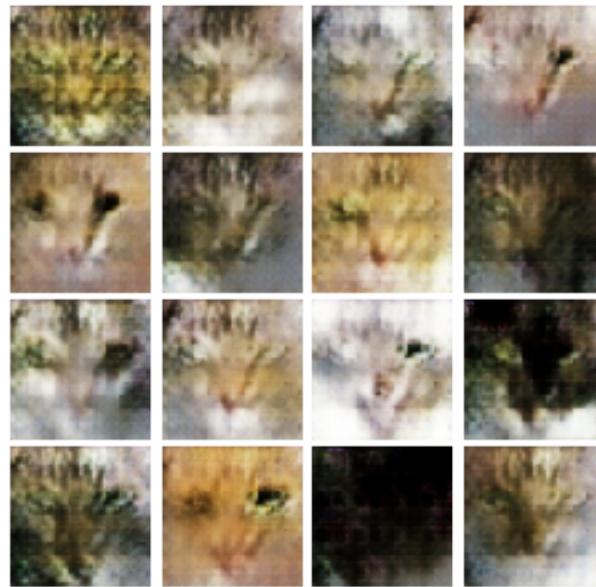
Iter: 1750, D: 0.189, G: 0.6023



EPOCH: 5  
Iter: 2000, D: 0.2069, G: 0.7338

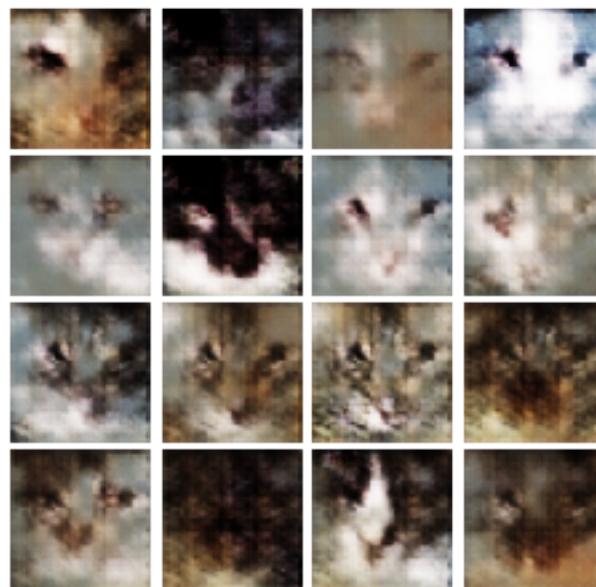


Iter: 2250, D: 0.5814, G: 0.6087

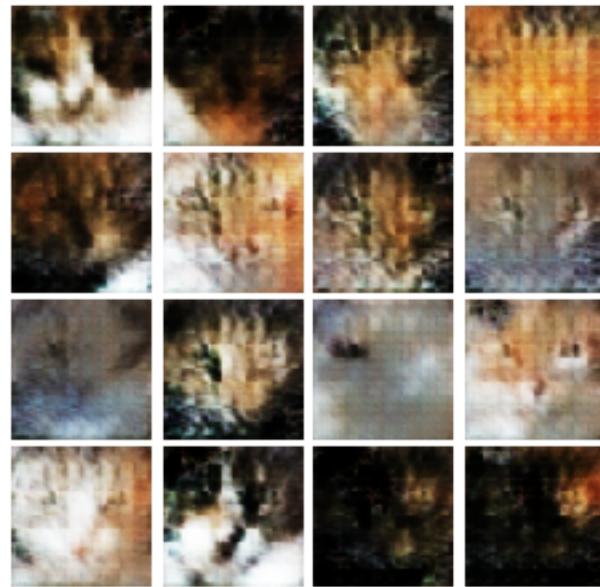


EPOCH: 6

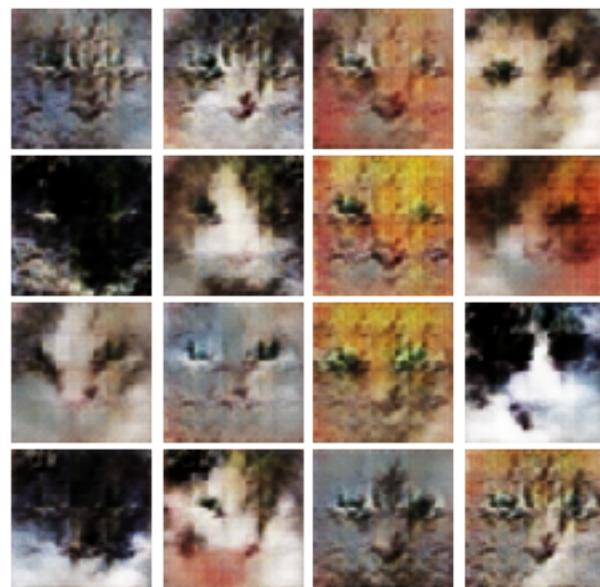
Iter: 2500, D: 0.1713, G:0.5729



Iter: 2750, D: 0.1876, G:0.764



EPOCH: 7  
Iter: 3000, D: 0.1538, G:0.9212



Iter: 3250, D: 0.1709, G:0.7127



EPOCH: 8  
Iter: 3500, D: 0.2482, G: 1.569



Iter: 3750, D: 0.06651, G: 1.194



EPOCH: 9

Iter: 4000, D: 0.1007, G:0.9568

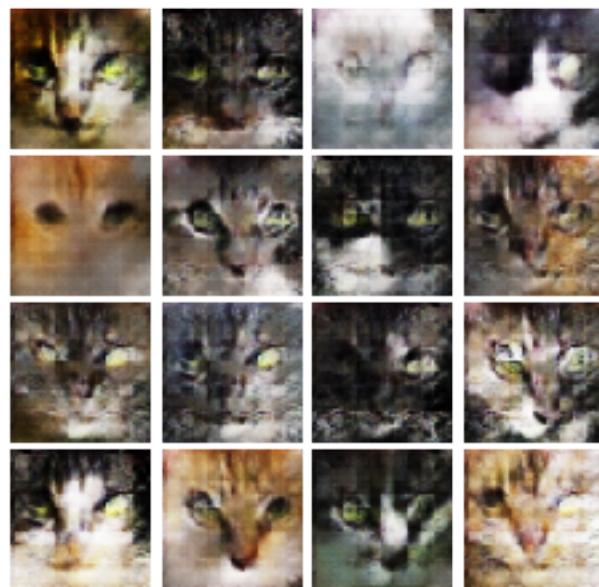


Iter: 4250, D: 0.0631, G:0.6403



EPOCH: 10

Iter: 4500, D: 0.09882, G:1.011



Iter: 4750, D: 0.1036, G:1.713



EPOCH: 11

Iter: 5000, D: 0.1462, G:1.594

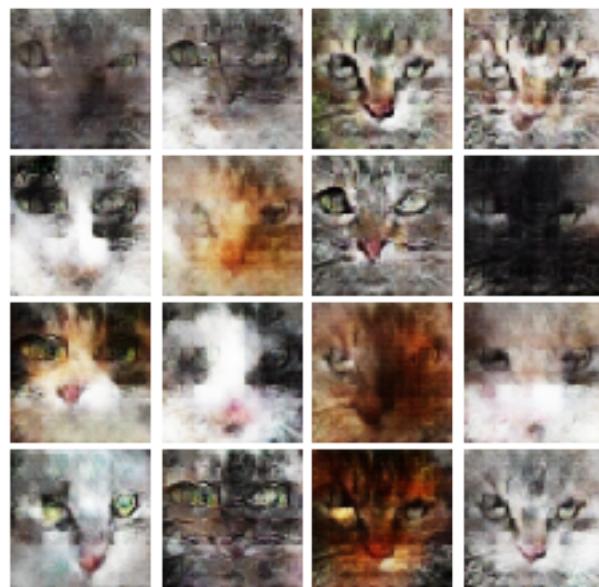


Iter: 5250, D: 0.02819, G:0.885



EPOCH: 12

Iter: 5500, D: 0.05475, G:0.9914



Iter: 5750, D: 0.04547, G:1.298



EPOCH: 13

Iter: 6000, D: 0.08845, G:1.26



Iter: 6250, D: 0.1224, G:1.739



EPOCH: 14

Iter: 6500, D: 0.06242, G:1.393



Iter: 6750, D: 0.05047, G:1.26



EPOCH: 15

Iter: 7000, D: 0.04122, G: 1.715



Iter: 7250, D: 0.05408, G: 1.429



EPOCH: 16

Iter: 7500, D: 0.04412, G: 1.012



Iter: 7750, D: 0.217, G: 0.3268



EPOCH: 17

Iter: 8000, D: 0.289, G: 0.7245



Iter: 8250, D: 0.07658, G: 1.425



EPOCH: 18

Iter: 8500, D: 0.09123, G: 0.8738



Iter: 8750, D: 0.0441, G: 0.8462



EPOCH: 19

Iter: 9000, D: 0.02897, G: 0.8615



Iter: 9250, D: 0.05496, G: 0.9372



EPOCH: 20

Iter: 9500, D: 0.05339, G:1.326



Iter: 9750, D: 0.01811, G:1.109

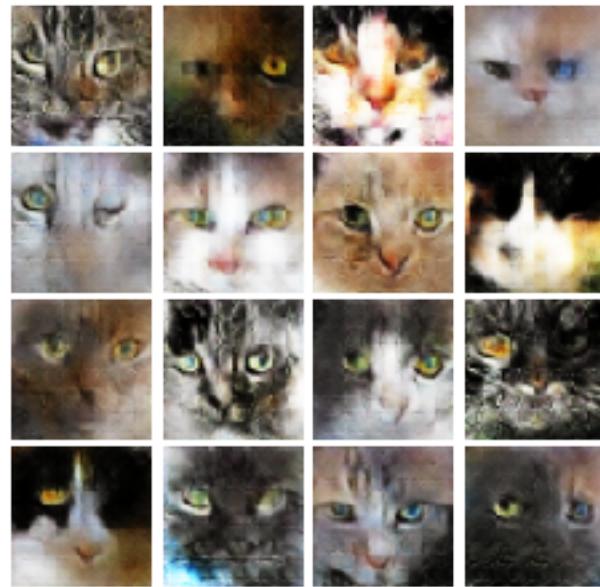


EPOCH: 21

Iter: 10000, D: 0.06138, G:1.144



Iter: 10250, D: 0.03175, G:0.5981



EPOCH: 22

Iter: 10500, D: 0.05124, G:1.549



Iter: 10750, D: 0.1603, G:0.4874



EPOCH: 23

Iter: 11000, D: 0.05563, G:1.103

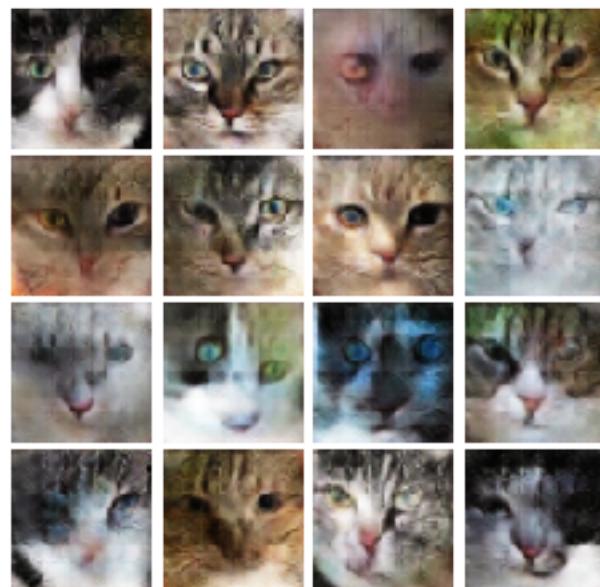


Iter: 11250, D: 0.1599, G:1.856



EPOCH: 24

Iter: 11500, D: 0.01807, G:1.109

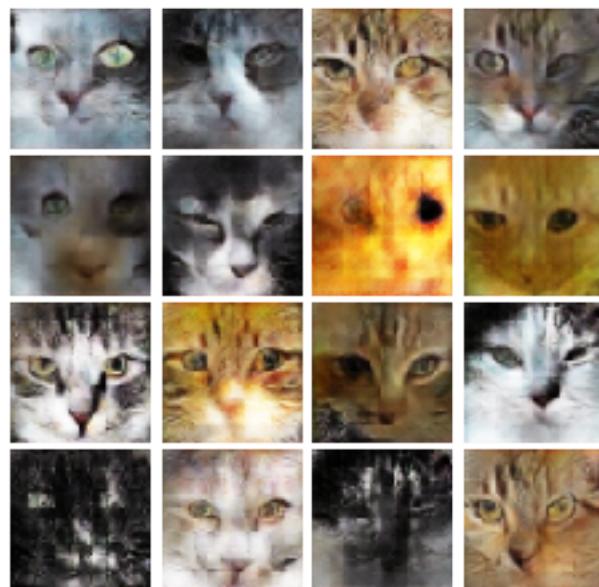


Iter: 11750, D: 0.0394, G:1.245



EPOCH: 25

Iter: 12000, D: 0.03951, G:0.848



Iter: 12250, D: 0.02886, G:0.6916



EPOCH: 26

Iter: 12500, D: 0.04181, G:0.917



Iter: 12750, D: 0.01627, G:1.242



EPOCH: 27

Iter: 13000, D: 0.0116, G:1.363

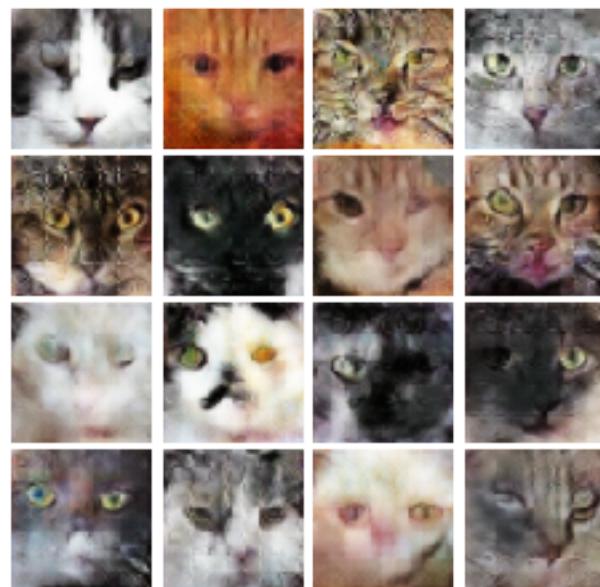


Iter: 13250, D: 0.03403, G:0.4368



EPOCH: 28

Iter: 13500, D: 0.04512, G:0.8367



Iter: 13750, D: 0.01979, G:0.9807



EPOCH: 29

Iter: 14000, D: 0.01616, G:1.115



Iter: 14250, D: 0.02699, G:1.219



EPOCH: 30

Iter: 14500, D: 0.1159, G:0.7933



Iter: 14750, D: 0.03234, G:0.9464



EPOCH: 31

Iter: 15000, D: 0.04441, G:0.7865

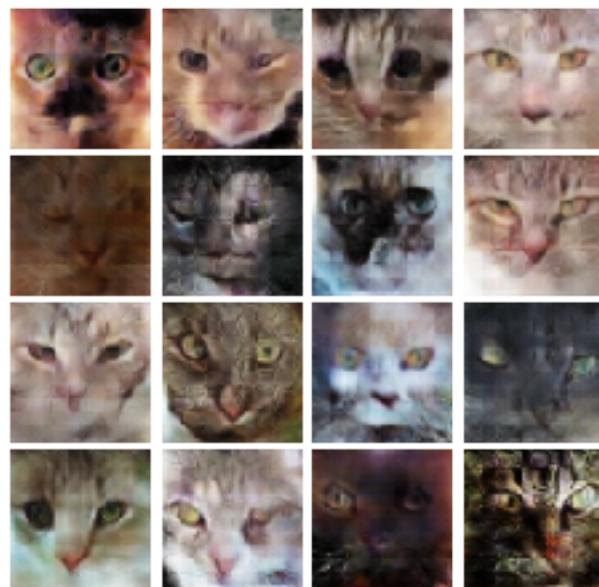


Iter: 15250, D: 0.07084, G:1.266



EPOCH: 32

Iter: 15500, D: 0.02577, G:1.219

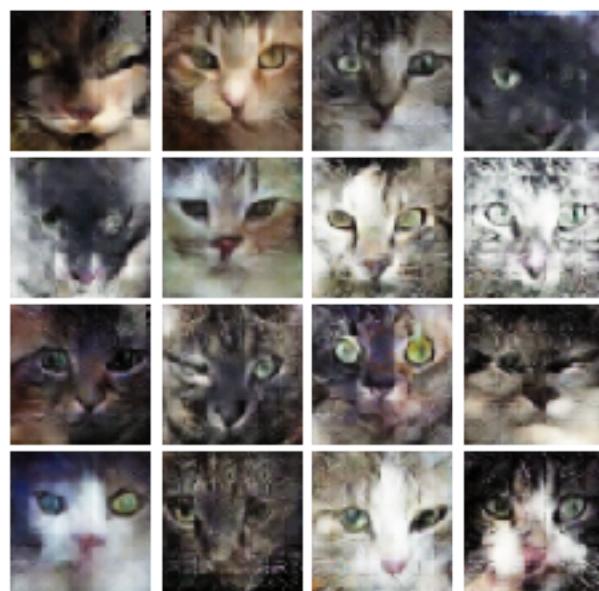


EPOCH: 33

Iter: 15750, D: 0.0295, G:1.529



Iter: 16000, D: 0.03307, G:1.14



EPOCH: 34  
Iter: 16250, D: 0.05605, G:0.7635



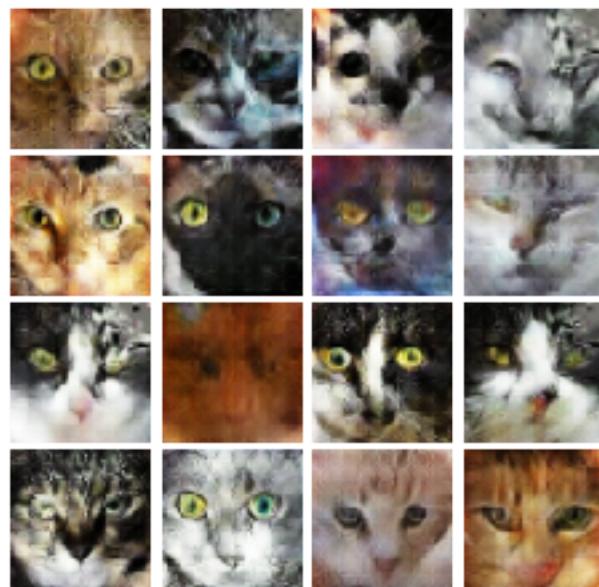
Iter: 16500, D: 0.02344, G:0.9316



EPOCH: 35  
Iter: 16750, D: 0.01, G:1.115

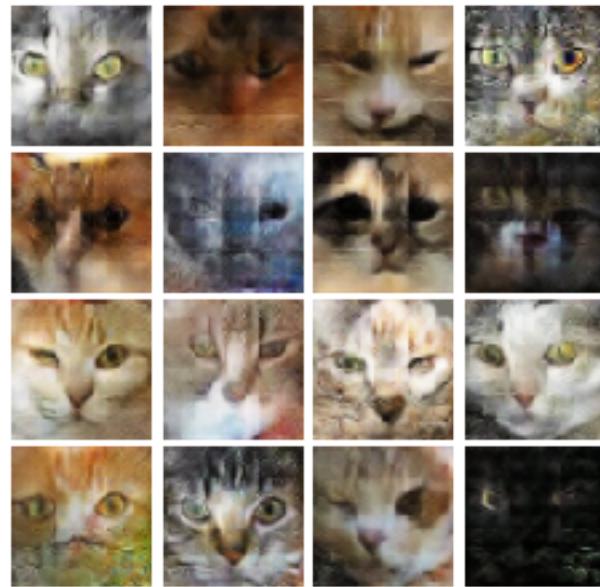


Iter: 17000, D: 0.01628, G:0.9987



EPOCH: 36

Iter: 17250, D: 0.04667, G:0.6505



Iter: 17500, D: 0.05236, G:0.5642



EPOCH: 37  
Iter: 17750, D: 0.02395, G:0.9202



Iter: 18000, D: 0.02232, G:1.035



EPOCH: 38  
Iter: 18250, D: 0.03833, G:0.7652



Iter: 18500, D: 0.01895, G:0.976



EPOCH: 39  
Iter: 18750, D: 0.0136, G:1.003



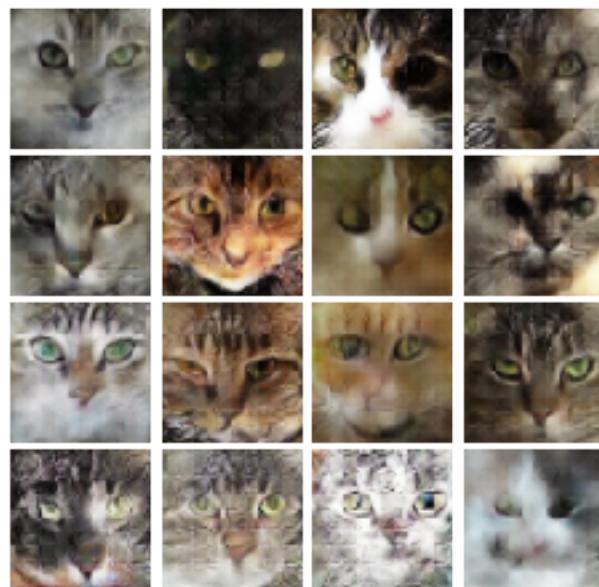
Iter: 19000, D: 0.01103, G:0.8844



EPOCH: 40  
Iter: 19250, D: 0.00873, G:0.9108



Iter: 19500, D: 0.05041, G:1.116



EPOCH: 41  
Iter: 19750, D: 0.06259, G:0.5558

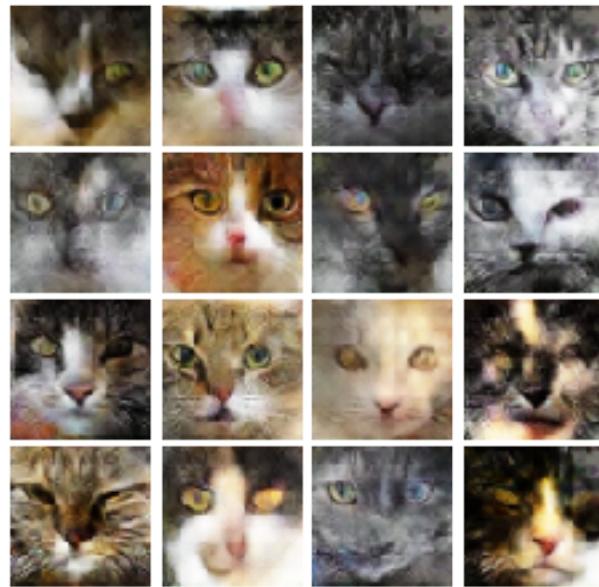


Iter: 20000, D: 0.02012, G:0.9161

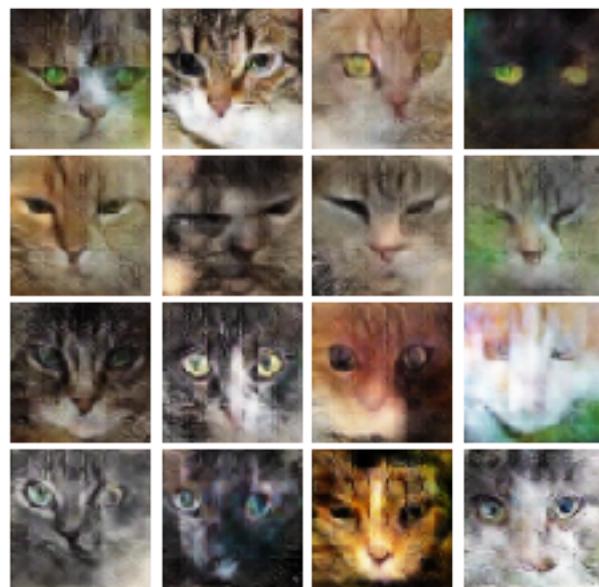


EPOCH: 42

Iter: 20250, D: 0.01332, G:0.8587



Iter: 20500, D: 0.0721, G:0.9979

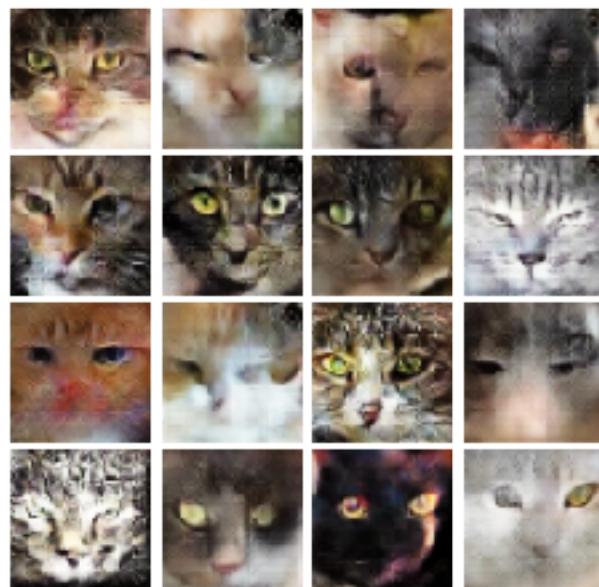


EPOCH: 43

Iter: 20750, D: 0.07577, G:0.7157



Iter: 21000, D: 0.1872, G:0.3597



EPOCH: 44

Iter: 21250, D: 0.0217, G:0.9086



Iter: 21500, D: 0.02811, G:1.044

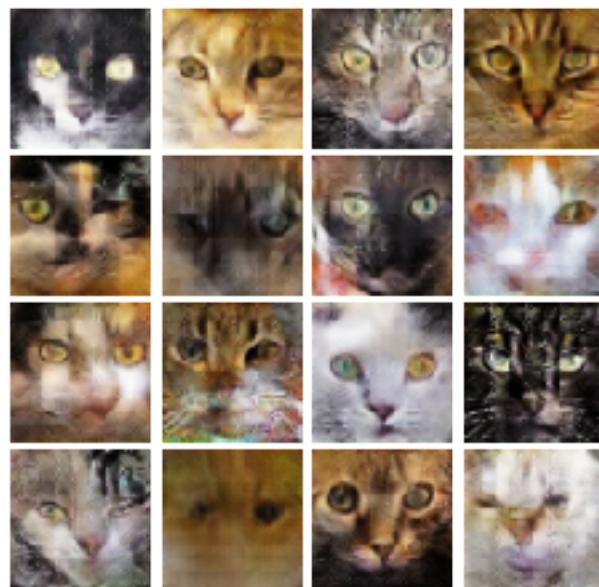


EPOCH: 45

Iter: 21750, D: 0.03301, G:1.239



Iter: 22000, D: 0.0119, G:0.9814

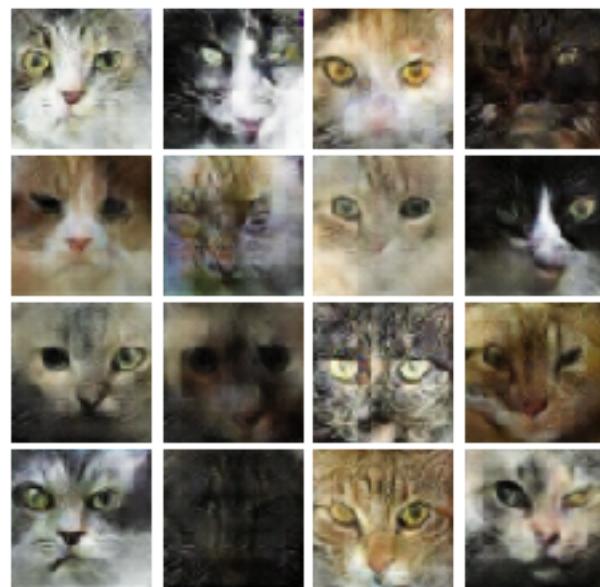


EPOCH: 46

Iter: 22250, D: 0.3501, G:0.6238



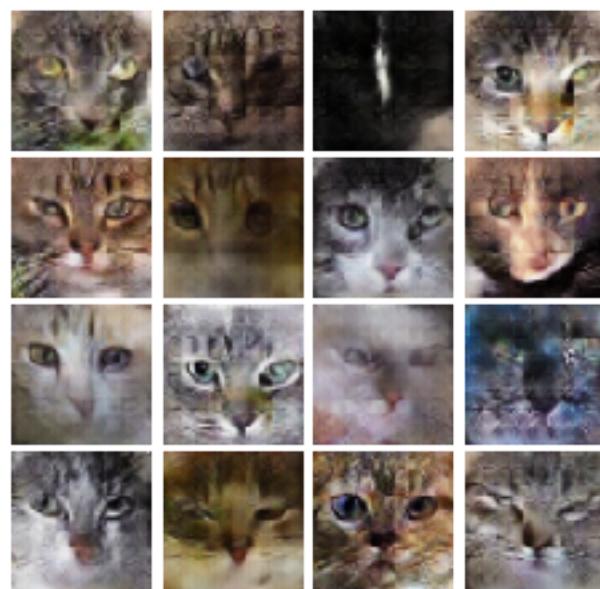
Iter: 22500, D: 0.00842, G:1.166



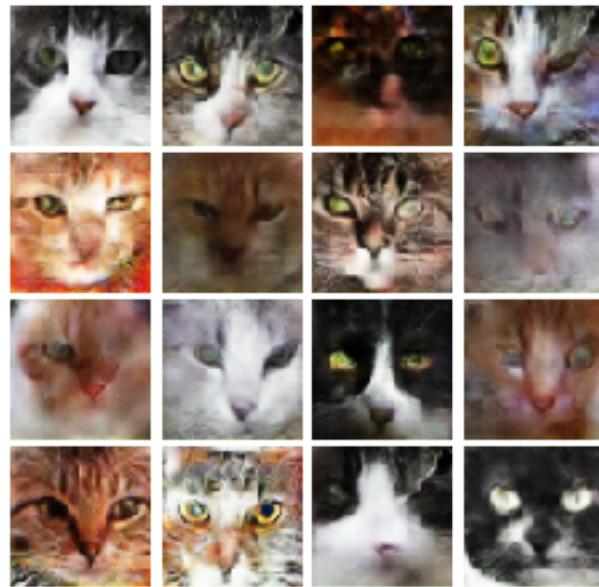
EPOCH: 47  
Iter: 22750, D: 0.01202, G:1.051



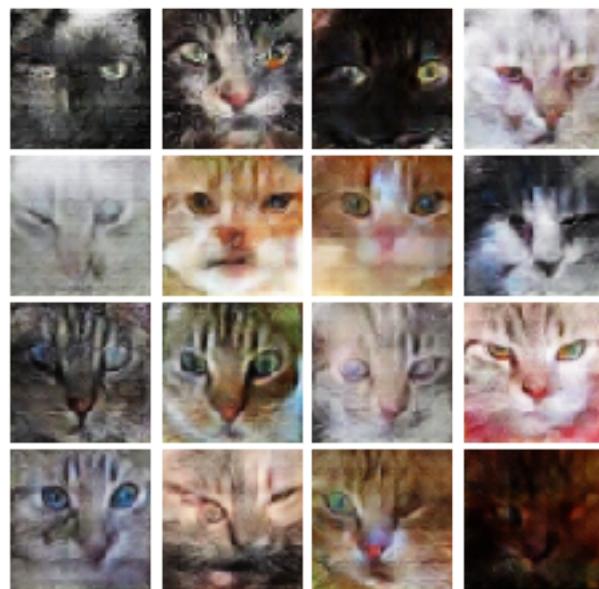
Iter: 23000, D: 0.02476, G:0.5607



EPOCH: 48  
Iter: 23250, D: 0.0204, G:1.52



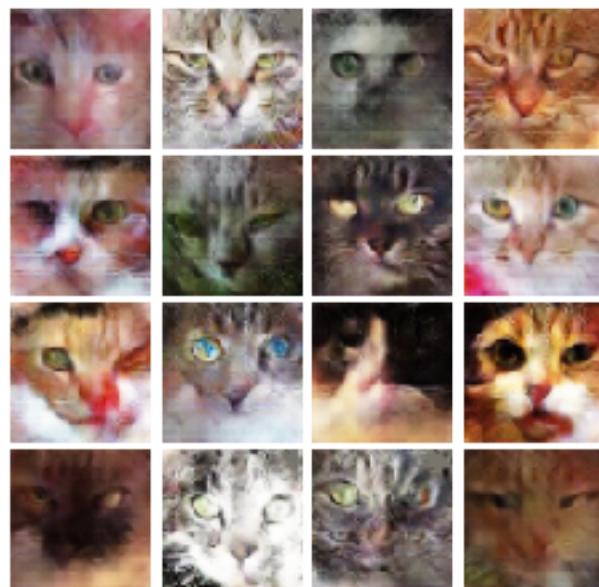
Iter: 23500, D: 0.04886, G:1.107



EPOCH: 49  
Iter: 23750, D: 0.02163, G:0.9345



Iter: 24000, D: 0.03091, G:1.112

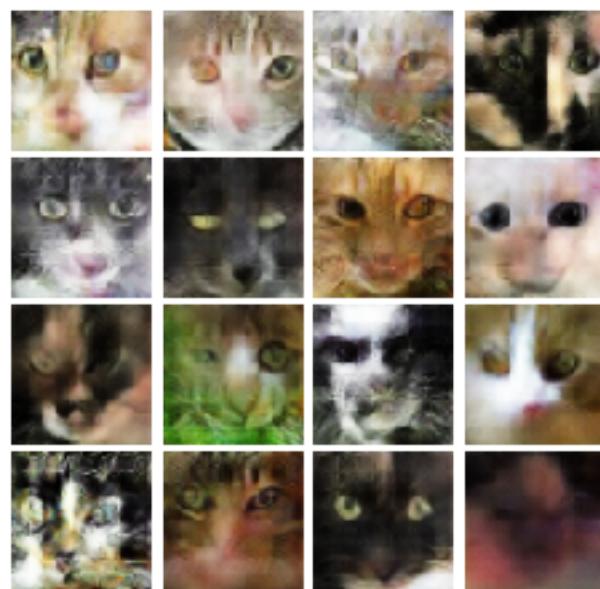


EPOCH: 50

Iter: 24250, D: 0.01475, G:1.392



Iter: 24500, D: 0.01059, G: 1.003



[ ]: